

# Implementation of Naive Bayes Classifier from Scratch

## Technical Report Presentation

Khamis K Haji (zda25m002) Juweayria Farouk (zda25m003)

Instructor: Dr. Innocent Nyalala

Course: Z5007: Programming and Data Structures

Institution: IIT Madras Zanzibar

January 19, 2026

## Why Implement from Scratch?

- Deepen understanding of Bayesian probability theory and classification
- Gain hands-on experience in designing efficient data structures
- Address real-world challenges:
  - Numerical underflow
  - Class imbalance
  - Scalability
- Benchmark against industry-standard tools (scikit-learn)
- Implement a custom hash table for storing model parameters

## Modular Design

### ① Custom Data Structures

- Hash table with dynamic resizing & collision handling

### ② Core Classifiers

- Gaussian, Bernoulli, Multinomial Naive Bayes

### ③ Preprocessing Pipeline

- Data loading, normalization, splitting

### ④ Testing Framework

- Unit, integration, performance tests

### ⑤ Evaluation Module

- Compares models using standard metrics

## Object-Oriented Design

- BaseClassifier (Abstract)
  - fit(X, y)
  - predict(X)
- GaussianNaiveBayes
  - \_mean\_table, \_variance\_table (HashTable)
- BernoulliNaiveBayes
  - \_likelihood\_table (HashTable)
- MultinomialNaiveBayes
  - feature\_log\_probs (ndarray)

## End-to-End Process

### ① Data Loading

- Wisconsin Breast Cancer dataset (569 samples, 30 features)

### ② Preprocessing

- Stratified train-test split (80:20)

### ③ Training

- Fit model using training data

### ④ Prediction

- Classify test samples

### ⑤ Evaluation

- Compare with scikit-learn baseline

# Custom Hash Table Implementation

## Key Features

- Separate chaining for collision resolution
- Dynamic resizing when load factor  $\geq 0.75$
- SHA-256 hashing for key distribution

## Complexity:

| Operation | Average Case | Worst Case |
|-----------|--------------|------------|
| Insert    | $O(1)$       | $O(n)$     |
| Search    | $O(1)$       | $O(n)$     |
| Resize    | $O(n)$       | $O(n)$     |

# Gaussian Naive Bayes

## Assumption & Training

- Assumes features follow Gaussian distribution:

$$x_j \sim \mathcal{N}(\mu_{yj}, \sigma_{yj}^2)$$

- Training:

$$\mu_{yj} = \frac{1}{n_y} \sum_{x \in y} x_j$$

$$\sigma_{yj}^2 = \frac{1}{n_y} \sum_{x \in y} (x_j - \mu_{yj})^2 + \epsilon$$

where  $\epsilon = 10^{-9}$

# Bernoulli & Multinomial Naive Bayes

## Bernoulli:

- Binary features:  $x_j \in \{0, 1\}$
- Training with Laplace smoothing ( $\alpha$ )

$$P(x_j = 1|y) = \frac{\text{count}(x_j = 1, y) + \alpha}{n_y + 2\alpha}$$

## Multinomial:

- Frequency count features
- Training in log-space to prevent underflow

$$\log P(x_j|y) = \log \left( \frac{\text{count}(x_j, y) + \alpha}{\sum_{k=1}^m \text{count}(x_k, y) + m\alpha} \right)$$

# Complexity Analysis

## Time & Space Complexity

| Operation               | GaussianNB             | Bernoulli/MultinomialNB |
|-------------------------|------------------------|-------------------------|
| Training                | $O(n \cdot m \cdot c)$ | $O(n \cdot m \cdot c)$  |
| Prediction (per sample) | $O(m \cdot c)$         | $O(m \cdot c)$          |

Table: Time complexity where  $n$  = samples,  $m$  = features,  $c$  = classes

Space Complexity:  $O(c \cdot m)$

# Experimental Results – Performance

## Wisconsin Breast Cancer Dataset

| Model                | Accuracy | Precision | Recall | F1-Score |
|----------------------|----------|-----------|--------|----------|
| GaussianNB (Ours)    | 96.46%   | 97.50%    | 92.86% | 95.16%   |
| GaussianNB (sklearn) | 97.37%   | 97.56%    | 93.02% | 95.22%   |
| BernoulliNB (Ours)   | 98.25%   | 98.25%    | 98.25% | 98.25%   |
| MultinomialNB (Ours) | 89.38%   | 89.47%    | 89.38% | 89.42%   |

Table: Performance comparison on Wisconsin Breast Cancer dataset

# Experimental Results – Speed

## Execution Time Comparison

| Operation               | Our Implementation | scikit-learn   | Ratio               |
|-------------------------|--------------------|----------------|---------------------|
| GaussianNB Training     | 0.0027s            | 0.0014s        | 1.93× slower        |
| BernoulliNB Training    | 0.0004s            | 0.00015s       | 2.67× slower        |
| MultinomialNB Training  | 0.002s             | 0.0023s        | 1.15× faster        |
| Prediction (per sample) | 0.00008s           | 0.00005s       | 1.60× slower        |
| <b>Average Training</b> | <b>0.0017s</b>     | <b>0.0013s</b> | <b>1.31× slower</b> |

Table: Speed comparison between implementations

# Key Findings & Ablation Study

## What We Learned

- **Accuracy:** Comparable to scikit-learn (within 0.91% for GaussianNB)
- **Speed:** Slightly slower due to Python-level loops
- **Stability:** Log-space calculations prevented underflow
- **Imbalance:** Stratified splitting ensured fair evaluation

## Ablation Insights:

- Removing log-space → underflow
- Removing variance smoothing → division-by-zero errors
- Without Laplace smoothing → poor performance on unseen features

# Conclusion & Future Work

## Successfully Demonstrated:

- ① Deepened theoretical understanding of Bayesian methods
- ② Built practical skills in data structure design
- ③ Achieved performance comparable to scikit-learn
- ④ Ensured reproducibility with documented code & math

## Future Work:

- Parallelize training using NumPy vectorization
- Extend to sparse data formats
- Implement online learning for streaming data
- Add support for categorical features

# References & GitHub Repository

## References

- ① Murphy, K. P. (2012). *Machine Learning: A Probabilistic Perspective*.
- ② Bishop, C. M. (2006). *Pattern Recognition and Machine Learning*.
- ③ Scikit-learn Naive Bayes Documentation
- ④ UCI Machine Learning Repository
- ⑤ Manning, C. D. et al. (2008). *Introduction to Information Retrieval*.

## GitHub Repository

Complete source code available at:

<https://github.com/fjuwearyia-dotcom/NaiveBayes>

# Thank You!

Questions?